

CHATTY STOCHASTIC MULTI-ARMED BANDITS

AKSHAY KUMAR

ADVISOR: SÉBASTIEN BUBECK

SUBMITTED IN PARTIAL FULFILLMENT

OF THE REQUIREMENTS FOR THE DEGREE OF

BACHELOR OF SCIENCE IN ENGINEERING

DEPARTMENT OF OPERATIONS RESEARCH AND FINANCIAL ENGINEERING

PRINCETON UNIVERSITY

JUNE 2014

I hereby declare that I am the sole author of this thesis.

I authorize Princeton University to lend this thesis to other institutions or individuals for the purpose of scholarly research.

Akshay Kumar

I further authorize Princeton University to reproduce this thesis by photocopying or by other means, in total or in part, at the request of other institutions or individuals for the purpose of scholarly research.

Akshay Kumar

Abstract

This thesis uses a variant of the classic stochastic multi-armed bandit framework to improve the user experience in an online chat application by selecting conversation starters. While the traditional algorithm would converge on the ‘optimal’ conversation starter and use it for every conversation, this novel version of the algorithm attempts to provide new conversation starters for each user while still attempting to maximize the conversation quality. This thesis examines the empirical behavior of such an algorithm in a web application deployed at Princeton University.

Acknowledgements

First and foremost, I would like to thank my parents and sister for years of love and support as I conclude my undergraduate education. I am the person that I am today because of them, and I am very proud to call them my family.

I am also grateful to Professor Bubeck for his guidance over the course of my senior year. His flexibility made this relatively un-orthodox ORFE thesis extremely enjoyable, and allowed me to spend my senior year getting excited about building something new. I am also grateful to my friends Zachary Garcia '14 and Damjan Korać '13 for giving excellent feedback on the first draft.

This thesis also wouldn't have been possible without the help of Daniel Kang '14, with whom I collaborated on the auxiliary, non-bandit related code in this thesis. His experience in web design and front-end development was indispensable in getting Tigers Anonymous up and running in such a short time, and his friendship made the collaboration that much more enjoyable.

Finally, I'd like to thank my wonderful girlfriend Isabella Chen '15 for being there for me throughout the course of this year. Through thick and thin, I knew she would always have my back.

To Mom, Dad, Neha and Isabella

Contents

Abstract	iii
Acknowledgements	iv
List of Figures	viii
1 Introduction	1
1.1 What is Tigers Anonymous?	1
1.2 Why Multi-Armed Bandits?	3
2 Background Research	4
3 Methods	8
3.1 UCB1-AKSB Algorithm	8
3.2 Tigers Anonymous Data Collection Methods	9
3.3 Tigers Anonymous Data Format	9
3.4 TA UCB1-AKSB Implementation	11
4 Results	14
4.1 Regret Analysis	15
4.2 UCB1-AKSB Effectiveness	17
4.3 Individual User Analysis	21
5 Conclusions	26

A Data Analysis Code	30
B Conversation Starters	36
C TA Back-End Implementation	40
C.1 Princeton IP-Address Filtering Functionality	40
C.2 User Matching Functionality	41
C.3 Web Server Functionality	45
C.4 Conversation Metadata Logging Model	47
D TA Front-End Implementation	48
D.1 Homepage	48
D.2 About Page	51
D.3 Chatroom	53

List of Figures

1.1	TA Conversation Starter	2
1.2	TA Drop-Down Menu	2
4.1	TA Google Analytics Dashboard	14
4.2	TA De-Anonymization Regret Analysis	15
4.3	TA Cumulative Conversation De-Anonymization Rate	17
4.4	TA Daily Conversation De-Anonymization Rate	19
4.5	TA Cumulative Participation Rate	20
4.6	TA Cumulative Average Conversation Length	20
4.7	De-Anonymization Rate Per User	21
4.8	Total Messages Sent Per User	22
4.9	Participation Rate Per User	22
4.10	User Engagement by Visit Count	24
4.11	User Engagement by Visit Length	24
D.1	Tigers Anonymous Homepage	48
D.2	Tigers Anonymous About Page	51
D.3	Tigers Anonymous Chatroom	53

Chapter 1

Introduction

At most college campuses, as students become more settled within their college community, it becomes increasingly difficult for them to branch out and meet new people beyond their immediate social graph. This is where TA comes in. By providing a way to chat anonymously with (and potentially meet) their fellow classmates, TA allows students to make new connections and shake up their social network. This website also provides a great way to develop and test a new, weakly context-dependent variant of the classic UCB1 multi-armed bandit algorithm (the UCB1-AKSB algorithm, described fully in Chapter 3). The development and performance of UCB1-AKSB is the main focus of this thesis.

1.1 What is Tigers Anonymous?

Tigers Anonymous (TA) is the name of a web application that allows any Princeton student to be matched with another Princeton student. After being matched, the students will be taken to an anonymous chatroom where they are given a conversation starter (an example is shown in Figure 1.1) and have the opportunity to have a conversation. For a complete list of all conversation starters used in TA, see Appendix B.

Once both participants have exchanged a pre-determined number of messages, a drop-down menu appears containing two choices (see Figure 1.2 below). If both users click ‘Yes’, the application will authenticate both users via Facebook and reveal each users’ identities to the other to facilitate communication outside of TA. For more information on how TA is implemented, see Appendix C and Appendix D.

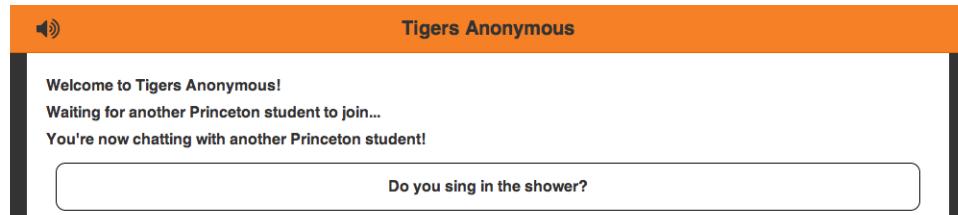


Figure 1.1: TA Conversation Starter

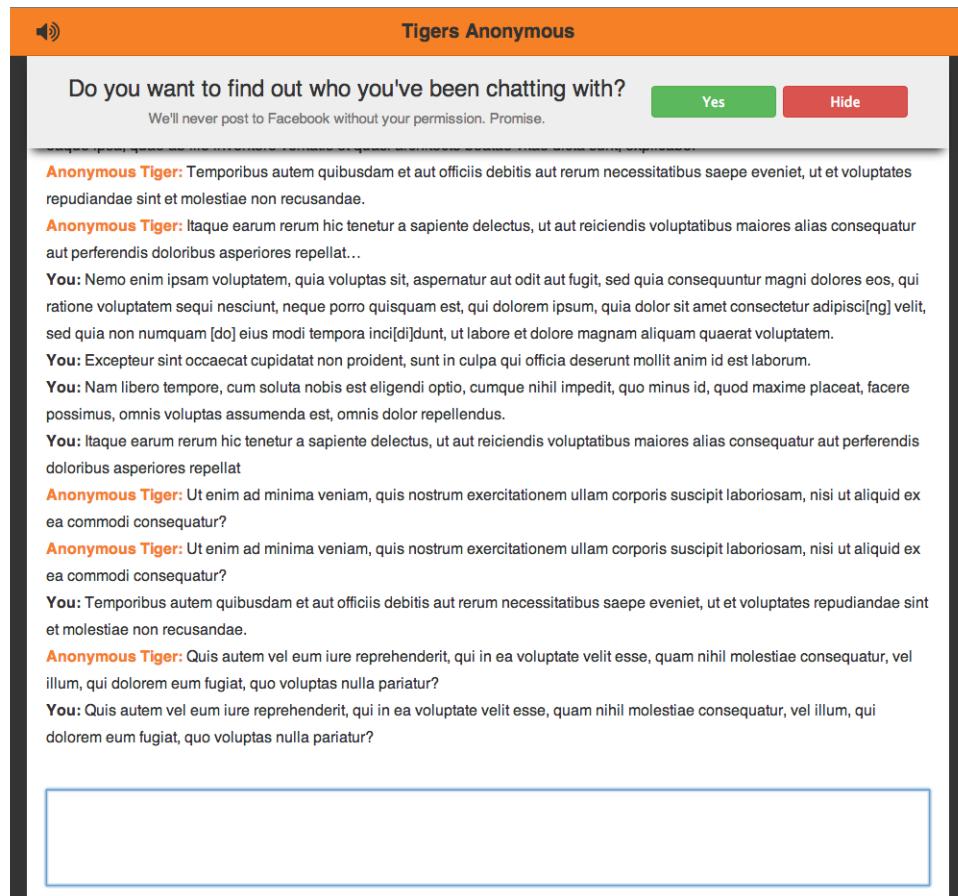


Figure 1.2: TA Drop-Down Menu

1.2 Why Multi-Armed Bandits?

A significant part of the functionality of TA is providing a conversation starter to reduce the awkwardness of the initial interaction with an anonymous stranger online.

A naive approach would simply choose conversation starters at random, but this approach would be less than optimal for two reasons. First, students might respond better to some conversation starters than others, so TA should be able to select the conversation starters that will facilitate higher quality conversations. Second, users could potentially see the same conversation starter more than once, which would defeat the purpose of having novel conversation starters in the first place.

This is where the multi-armed bandit problem comes in. By modeling conversation starters as ‘arms’ in a classical multi-armed bandit problem and a ‘success’ as a ‘high-quality’ conversation, it should be possible to solve both of the above problems. For more information on the multi-armed bandit problem and the motivation for the UCB1-AKSB algorithm described in Chapter 3, see Chapter 2.

Chapter 2

Background Research

In its most general form, the multi-armed bandit problem is a sequential decision-making problem where the decision-maker must explore new possibilities while trying to exploit existing knowledge in order to maximize the received reward. Using the notation in Bubeck and Cesa-Bianchi (2012), there are are $K \geq 2$ arms and sequences $X_{i,1}, X_{i,2}, \dots$ for each arm $i = 1, \dots, K$. At each time step $t = 1, 2, \dots$ the decision-making algorithm picks $I_t \in \{1, \dots, K\}$ and receives a reward $X_{I_t,t}$.

The common way of benchmarking the performance of a such an algorithm is to take the distance between the algorithm's expected reward and the optimal reward by play n . In the literature, this distance can be defined at two levels of granularity: at the play-by-play level or at the arm-level. Using the first notion of optimality yields the expected regret, which is defined in Equation 2.1.

$$\mathbb{E}R_n = \mathbb{E} \left[\max_{i=1, \dots, K} \sum_{t=1}^n X_{i,t} - \sum_{t=1}^n X_{I_t,t} \right] \quad (2.1)$$

A slightly weaker notion of regret comes from the second definition of optimality: picking the arm with the optimal expected reward. Using this notion of optimality yields the pseudo-regret, which is defined in Equation 2.2.

$$\bar{R}_n = \max_{i=1,\dots,K} \mathbb{E} \left[\sum_{t=1}^n X_{i,t} - \sum_{t=1}^n X_{I_t,t} \right] \quad (2.2)$$

It is important to notice here that in Equation 2.1, the algorithm is competing with the best possible arm at every time-step, whereas in Equation 2.2, the algorithm is only competing with the arm with the highest expected value. This is the intuition for why $\bar{R}_n \leq \mathbb{E} R_n$.

Building upon this general problem formulation, there are three main sub-problems outlined by Bubeck and Cesa-Bianchi (2012) with further assumptions about the reward distributions of the bandit arms. In the stochastic version, the sequences of plays $X_{i,1}, X_{i,2}, \dots$ are I.I.D samples drawn from distributions $\nu_i \in [0, 1]$ for each arm $i = 1, \dots, K$. In the adversarial version, at each time step t an adversary selects a gain vector $g_t = (g_{1,t}, \dots, g_{K,t}) \in [0, 1]^K$ such that $X_{i,t} = g_{i,t}$. In the Markovian version, the reward processes are neither i.i.d (as in the stochastic version) or chosen by an adversary (as in the adversarial version). Instead, each arm represents a Markov process with a state space S . Each time an arm i is chosen in state $s \in S$, a reward is drawn from a probability distribution $\nu_{i,s}$ and the state of the Markov process for arm i changes to $s' \in S$ based on a transition probability matrix M_i .

Armed with these three canonical variants of the multi-armed bandit problem, we are ready to tackle the problem at hand: choosing conversation starters for a pair of TA users. Assuming that Princeton students will react similarly to different conversation starters, it seems intuitive to use the stochastic formulation of the multi-armed bandit problem since the reward distributions for each arm could reasonably be assumed to be I.I.D. Additionally, the stochastic bandit problem has a well-documented solution: the UCB1 algorithm, which is elegant, computationally easy to implement, and has a uniformly logarithmic upper-bound on cumulative pseudo-regret (Auer, Cesa-Bianchi, and Fischer, 2002).

Although modeling the problem in this way solves the first problem outlined in

Section 1.2 (i.e. picking the conversation starter that students would respond best to), it would actually exacerbate the second problem (i.e. seeing the same conversation starter repeatedly). Since the stochastic multi-armed bandit version of the problem assumes the arm reward distributions are I.I.D, it will hone in on the ‘best’ arm, and in the long term, will simply display the same conversation starter over and over again. This, in turn, would defeat the purpose of having a variety of conversation starters in the first place.

Since the arms chosen need to depend on which users are chatting, a logical first choice would be to turn to a contextual bandit algorithm. One heavily-cited example of such an algorithm (LinUCB) can be found in Li, Chu, Langford, and Schapire (2010). Li et al. (2010) observe the current user and a set of arms along with a feature vector for each user/arm pair. This vector summarizes information of both the user and arm, and can be thought of as the context. However, such a model would be infeasible for TA for a variety of reasons. First, since the context is defined by a pair of users, one would have to maintain an extremely sparse database of all feature vectors for all arms, for all possible pairs of users. With a reasonable user-base of 1000 users and 200 conversation starters, this would result in 200,000,000 (e.g. $200 * 1000 * 1000$) feature vectors, many of which would be empty or very sparse. Second, the actual algorithm proposed in Li et al. (2010) requires matrix multiplication and inversion to calculate the UCB value for each contextual bandit, which makes it computationally costly and difficult to implement practically for a web application handling multiple concurrent requests on a single server. Finally, having such a high level of context-dependency is unnecessary given the assumption that Princeton students’ responses to a given conversation starter will be I.I.D.

This raises an important question: how do we take advantage of the simplifying assumption that Princeton students will respond similarly to any given conversation starter (i.e. the rewards are I.I.D) while still enforcing the context-dependent

invariant that a user doesn't see the same conversation starters repeatedly? This was the motivation to create a UCB algorithm somewhere between the context-free UCB1 and the context-dependent LinUCB. This new algorithm (which I have named UCB1-AKSB) is weakly context dependent, in that it applies the context-free UCB1 algorithm but dynamically filters the arms over which the algorithm to only include arms that neither user has already seen. This intuitive explanation is formalized in Chapter 3 below.

Chapter 3

Methods

3.1 UCB1-AKSB Algorithm

The multi-armed bandit algorithm used by Tigers Anonymous is a novel variant of the well-known UCB1 algorithm (Auer et al., 2002). The new algorithm is outlined below:

Before explaining the algorithm, it will be useful to introduce notation. Let the users be represented as the set U and the bandit arms as the set X . Let the set of arms that have already been played for user $u \in U$ be represented by the set $X^u \subset X$. The goal of the UCB1-AKSB algorithm is to pick some arm $x \in X$ given the pair of users $u, v \in U$. In this specific application, the goal is to pick the optimal conversation starter $x \in X$.

The UCB1-AKSB algorithm proceeds as follows: For each pair of users $u, v \in U$, we pick the conversation starter x such that

$$x = \arg \max_{x \in \{X^u \cup X^v\}^c} f(x) \quad (3.1)$$

where

$$f(x) = \begin{cases} \bar{x} + \sqrt{\frac{2\ln n}{n_x}} & : n_x > 0 \\ \infty & : n_x = 0 \end{cases} \quad (3.2)$$

In Equation 3.2, n_x is the number of times that conversation starter x has been played and n is the total number of conversation starters that have been shown. Note that ties are broken arbitrarily. Additionally, in the event that $\{X^u \cup X^v\} \in \emptyset$, the algorithm simply selects a random arm.

3.2 Tigers Anonymous Data Collection Methods

The complete data-collection method used for this thesis is outlined below:

1. Two users visit www.tigersanonymous.com/chat from a Princeton IP address.
2. The users are directed to the chat server and are matched on a first-come, first-served basis.
3. A conversation starter is selected based on the UCB1-AKSB algorithm described above.
4. After either of the users disconnects, a 10-tuple representing the chat session is logged in a database (see Data Format section below for more details).

3.3 Tigers Anonymous Data Format

The data that will be collected can be represented by the vector of 10-tuples of the form:

$$(x_i, y_i, t0_i, t1_i, q_i, b_i, c1_i, c2_i, m1_i, m2_i)$$

where x_i and y_i represent the pseudonymous user ids of the two participants in the chat, $t0_i$ and $t1_i$ represent the start and end times of the conversation, q_i represents

the conversation starter, $b_i \in (0, 1)$ represents whether the drop-down menu was displayed (i.e. both chat participants exchanged more than a predefined number of messages), $c1_i, c2_i \in (0, 1)$ represent whether users x_i and y_i opted to de-anonymize the conversation respectively and $m1_i, m2_i \in (0, 1)$ represent the number of messages that user x_i and y_i sent respectively. The subscript i is unique for each conversation.

A sample of this data is shown below:

```
[
  {
    "userID1" : "a8262bb13e641e2bf5dcb3985b2061be",
    "userID2" : "9a675a6f581fd1dfa0b982826e75b4f5",
    "question" : "Do you believe in soul mates?",
    "startTime" : 1390873219878,
    "endTime" : 1390873263469,
    "buttonDisplayed" : false,
    "user1Clicked" : false,
    "user2Clicked" : false,
    "user1MessagesSent" : 1,
    "user2MessagesSent" : 2,
    "_id" : "52e70aafc43b6d020079e52e",
    "__v" : 0
  },
  {
    "userID1" : "370f85e443ad3ee24a879b1ce5a2b54b",
    "userID2" : "6e0fe76fca80cf2920bd5fc7717cf6dd",
    "question" : "What's one thing that you learned this week?",
    "startTime" : 1390881063228,
    "endTime" : 1390882681992,
    "buttonDisplayed" : true,
    "user1Clicked" : true,
    "user2Clicked" : true,
    "user1MessagesSent" : 33,
    "user2MessagesSent" : 37,
    "_id" : "52e72f79c43b6d020079e531",
    "__v" : 0
  },
  ...
]
```

3.4 TA UCB1-AKSB Implementation

This is the code on the Tigers Anonymous chat server that implements the UCB1-AKSB algorithm.

```
var questions = require('./questions').list;

// Used in lieu of positive and negative infinity
var largePositiveNumber = 1000000000;
var largeNegativeNumber = -1000000000;

// UCB1-AKSB algorithm to pick conversation starter
exports.getQuestion = function(collection, user1, user2, callback) {
    var questionAsked = {
        $or: [
            {$eq: ["$userID1", user1.id]},
            {$eq: ["$userID2", user1.id]},
            {$eq: ["$userID1", user2.id]},
            {$eq: ["$userID2", user2.id]}
        ]
    };
}

var outputFormat = {
    _id: "$question",
    plays: {$sum: 1},
    wins: {$sum: {$cond: [{$and: ["$user1Clicked", "$user2Clicked"]}, 1, 0]}},
    timesShown: {$sum: {$cond: [questionAsked, 1, 0]}}
};

// Aggregate conversation data and invoke UCB1 on set of
// never-before-seen arms
collection.aggregate().group(outputFormat).exec(function(err, data) {
    if (err) console.log(err);
    UCB1(data, callback);
});

// Helper function to get a random question
var getRandomQuestion = function() {
    var randomIndex = Math.floor(Math.random() * questions.length);
    return questions[randomIndex];
};

// Calculate UCB1 values for all possible arms in data
```

```

var UCB1 = function(data, callback) {
  var finalData = {};

  // If there's no data, return a random question
  if (data.length === 0) {
    callback(getRandomQuestion());
    return;
  } else {
    // Otherwise, get all the available data for the questions and run UCB
    var questionStats = {};
    var totalPlays = 0;

    // For each entry in data, sum the total number of plays and
    // populate the questionStats table with the corresponding question
    for (var i = 0; i < data.length; i++) {
      var entry = data[i];
      questionStats[entry._id] = {
        plays: entry.plays,
        wins: entry.wins,
        shown: (entry.timesShown > 0 ? true : false)
      };
      totalPlays += entry.plays;
    }

    for (var i = 0; i < questions.length; i++) {
      var question = questions[i];
      // If there's no data for this question, then it hasn't been
      // displayed yet, so assign it an arbitrarily large UCB value
      if (!questionStats[question]) {
        finalData[question] = largePositiveNumber;
      } else if (questionStats[question].shown) {
        continue;
      } else {
        // If the question hasn't been shown and there's data for it,
        // compute the UCB value
        var probabilityEstimate =
          questionStats[question].wins / questionStats[question].plays;
        var UCBoundEstimate =
          Math.sqrt(2 * Math.log(totalPlays /
            questionStats[question].plays));
        finalData[question] = probabilityEstimate + UCBoundEstimate;
      }
    }

    if (Object.keys(finalData).length > 0) {
      // Find question with max UCB1 value
      var bestValue = largeNegativeNumber;

```

```
var bestMatch = null;
for (var question in finalData) {
    var currentValue = finalData[question];
    if (currentValue >= bestValue) {
        bestMatch = question;
        bestValue = currentValue;
    }
}
callback(bestMatch);
} else {
    callback(getRandomQuestion());
}
}
};


```

Chapter 4

Results

Overall, using TA to collect data for the UCB1-AKSB algorithm was a success. Not only did the site facilitate over 8,000 conversations between Princeton students that would have otherwise never occurred, but the site also gained significant traction on campus. As evidenced by the Google Analytics dashboard shown in Figure 4.1, TA had over 40,000 page-views and over 6,000 unique visitors. Although site activity stabilized at only modest levels, this was probably due to user stratification and/or saturation, which is explored in further detail in Section 4.3.



Figure 4.1: TA Google Analytics Dashboard

4.1 Regret Analysis

Recall from Chapter 2 that the cumulative regret of the UCB1 algorithm is proportional to $\log(n)$, where n is the number of plays. Thus, we should expect the UCB1-AKSB algorithm to result in a cumulative pseudo-regret that is approximately logarithmic in the number of plays. This empirical pseudo-regret, \hat{R}_n , was calculated using Equation 4.1 below under the assumptions that the long-term average de-anonymization rate (μ^*) was optimal.

$$\hat{R}_n = \mu^* n - \sum_{t=1}^n X_{I_t, t} \quad (4.1)$$

In Equation 4.1, $\mu^* n$ is the expected optimal number of de-anonymizations by play n (i.e. $\max_{i=1,\dots,K} \mathbb{E}[\sum_{t=1}^n X_{i,t}]$ from Equation 2.2). The plot of \hat{R}_n as a function of n for the UCB1-AKSB algorithm is shown below in Figure 4.2.

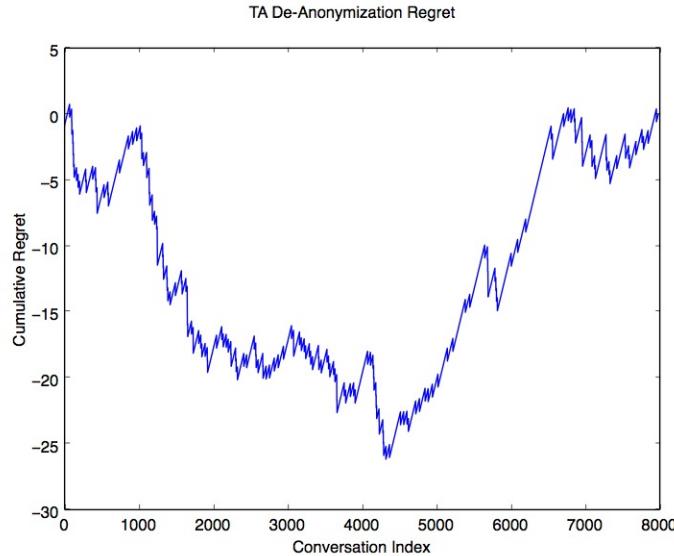


Figure 4.2: TA De-Anonymization Regret Analysis

The regret looks approximately logarithmic for the second half of the dataset, but the first half of the data gives a steadily negative cumulative regret. This is because the assumption that conversation de-anonymization is I.I.D is most likely false.

Instead, there were probably different regimes in which people perceived Facebook de-anonymization differently. This is most analogous to the Markovian bandits in Bubeck and Cesa-Bianchi (2012), where each conversation starter is associated with a Markov process with a discrete set of reward distributions.

Because of the clear split in the data, it seems like there were two discrete distributions from which de-anonymizations were drawn. The first distribution occurred in the initial stages of TA’s launch, where users were more likely to have long conversations and de-anonymize the conversation simply because of the novelty of doing so. This is supported by looking at the initial cumulative Facebook de-anonymization statistics (see Figure 4.3), where the Facebook connect rate was almost double the long-term average. The second distribution most likely occurred as the novelty of de-anonymization wore off and users opted for de-anonymization only if the conversation was of sufficiently high quality. The existence of multiple regimes explains the two parts of the regret data in Figure 4.2.

However, the second part of the cumulative regret plot in Figure 4.2 still looks logarithmic, which suggests that the algorithm performed in-line with expectations during the second regime of TA usage.

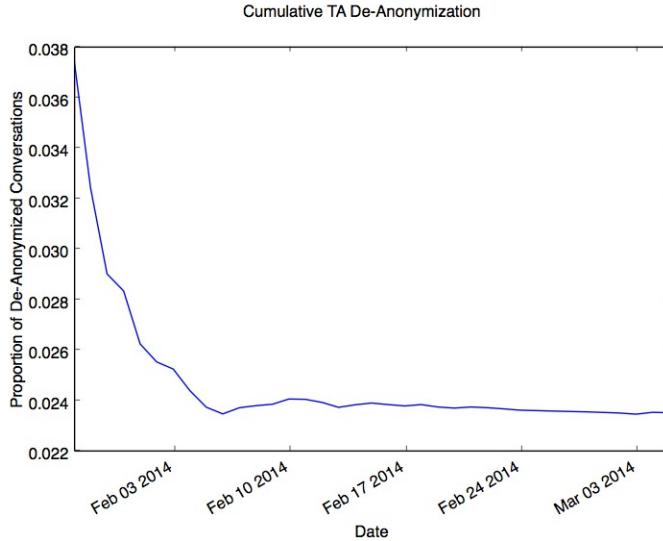


Figure 4.3: TA Cumulative Conversation De-Anonymization Rate

4.2 UCB1-AKSB Effectiveness

Since conversation de-anonymization was the metric that UCB1-AKSB used to judge each conversation starter, the obvious way to measure the performance of the UCB1-AKSB algorithm is the proportion of conversations which were de-anonymized. Judging from not only the cumulative de-anonymization proportion (Figure 4.3), but also the daily de-anonymization proportion (Figure 4.4), it seems that the algorithm had little impact on whether or not people opted to de-anonymize the conversation. If anything, it looks like the algorithm had an adverse impact on conversation de-anonymization.

However, this may have been due to the fact that the user's decision to de-anonymize the conversation was based on factors other than the conversation starter. It is easy to imagine a case in which the conversation was extremely revealing and thus users were hesitant to reveal their identities for fear of being connected to the conversation. In such cases, the conversation starter may have been excellent

but the conversational and/or social tendencies of the users would prevent them from de-anonymizing. In other words, the process of conversation de-anonymization might have been more user-specific (i.e. context dependent) than the UCB1-AKSB algorithm assumed, which could have resulted in a downward drag on overall de-anonymization rates even as the conversation starter quality improved.

Additionally, recall that the option to de-anonymize a conversation only came after both users had exchanged a pre-specified number of messages (in the current implementation, this number is 15). The implication is clear: it is possible that many of these conversations never even had the opportunity to de-anonymize, so marking them as ‘failures’ in the data model is misleading. This claim is empirically validated by aggregate TA user data as shown in ???. This graph shows that, over time, users became more likely to de-anonymize the conversation, which suggests that the UCB1-AKSB algorithm was improving the quality of conversation starters despite the overall drop in de-anonymizations. In other words, even though the proportion of people who de-anonymized decreased slightly, the proportion of people who de-anonymized when given the opportunity actually increased, which suggests that UCB1-AKSB may have been at least moderately effective in spurring conversation de-anonymization.

In any case, the metric of conversation de-anonymization itself introduces difficulties in accurately measuring incremental improvements in conversation quality. For example, let us say the conversation de-anonymization would only occur if the conversation ‘quality’ metric was above some threshold d . Even if the UCB1-AKSB algorithm boosted the quality of otherwise low-quality conversation by providing some common ground, such a quality boost would not be visible unless the incremental improvement was enough to make such conversations pass the threshold d . Basically, it is entirely possible that the UCB1-AKSB algorithm could improve conversation quality but this increased likelihood would not be visible because of the censored data observed. This ‘censored-data’ hypothesis would also explain some of the erratic

behavior of the daily de-anonymization rate seen in Figure 4.4.

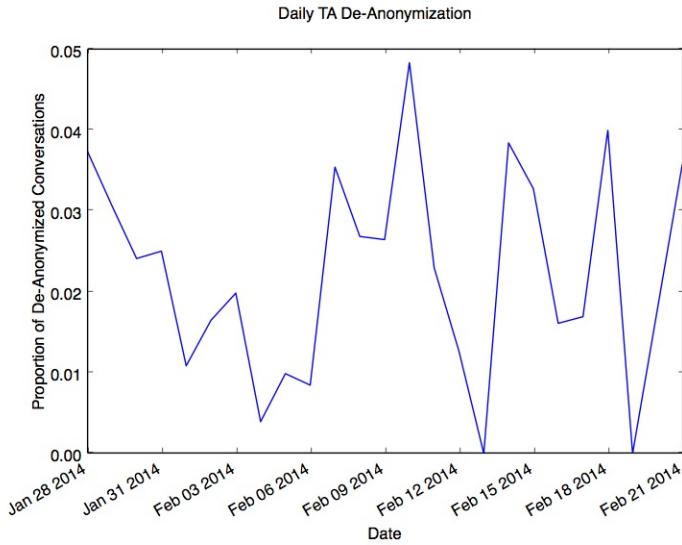


Figure 4.4: TA Daily Conversation De-Anonymization Rate

Given that conversation de-anonymization might have been affected by other exogenous variables and was not granular enough to measure incremental improvements in conversation quality, it makes sense to turn to other conversation quality metrics to judge the performance of UCB1-AKS. These other metrics (participation rates and average conversation rates) are less likely to be influenced by the social pressure for or against de-anonymization, as well as having a more finely differentiated set of values than the binary variable of conversation de-anonymization. The plots of both these metrics are shown below in Figure 4.5 and Figure 4.6.

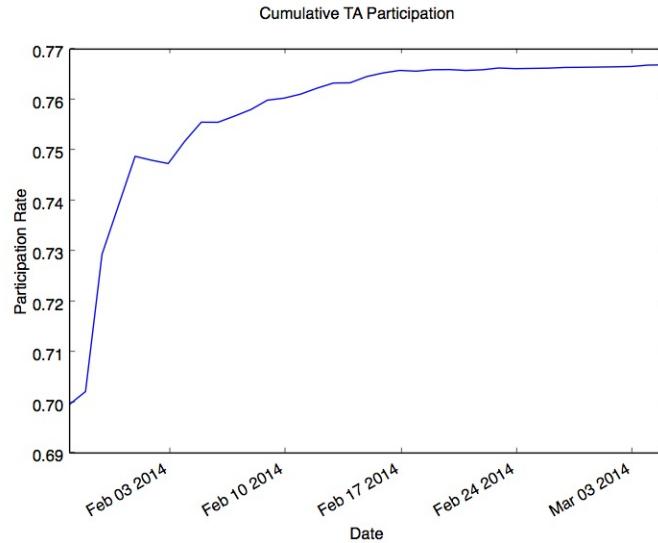


Figure 4.5: TA Cumulative Participation Rate

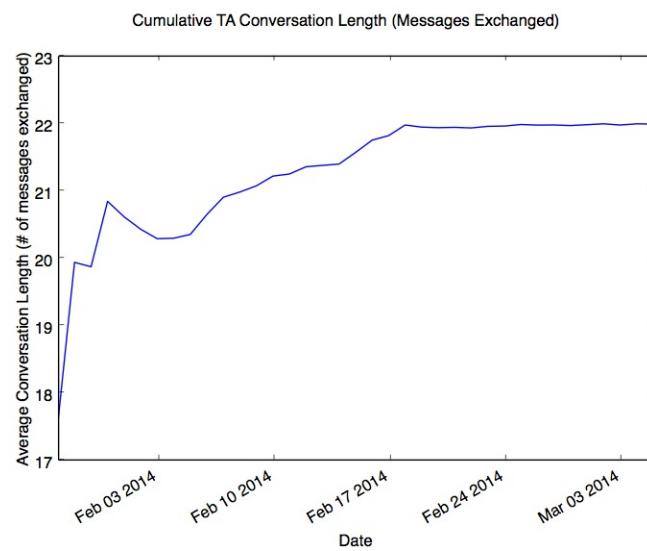


Figure 4.6: TA Cumulative Average Conversation Length

By these metrics, it seems that conversation quality improved noticeably over time, which suggests that the UCB1-AKSB algorithm may still have had a positive effect on conversation quality even though some of its fundamental assumptions were not true.

4.3 Individual User Analysis

Another way of examining the data is to look at how individual users behaved as they continued to interact with the site. In each of the plots below, the x-axis represents the number of uses, while the y-axis represents the conditional mean of the metric over the set of users on the n -th use, given that they have used the site at least n times. In order to define this more clearly, I introduce the following notation: let $f_k(u, n)$ give the value of conversational quality metric k for user u on their n -th visit and function $g(u)$ give the number of times user u has visited the site. Let U_n be the set $\{u | u \in U, g(u) \geq n\}$ (i.e. the set of users who have visited the site at least n times). Then, the graphs below are plots of the following functions $y_k(n)$ for different conversational quality metrics k .

$$y_k(n) = \frac{1}{|U_n|} \sum_{u \in U_n} f_k(u, n) \quad (4.2)$$

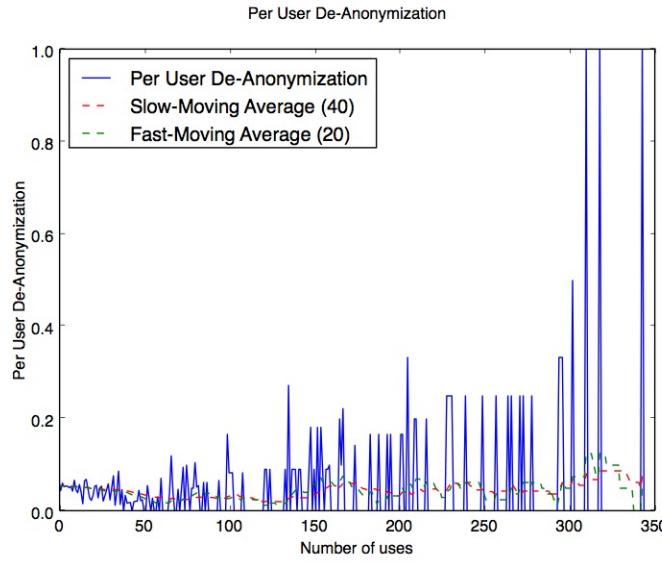


Figure 4.7: De-Anonymization Rate Per User

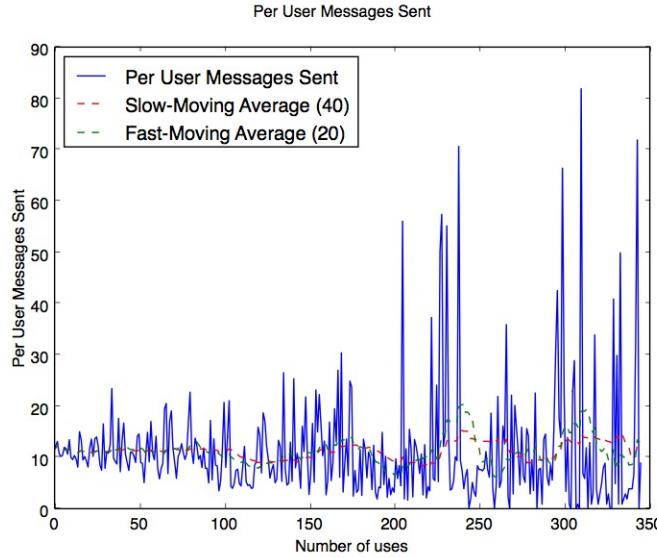


Figure 4.8: Total Messages Sent Per User

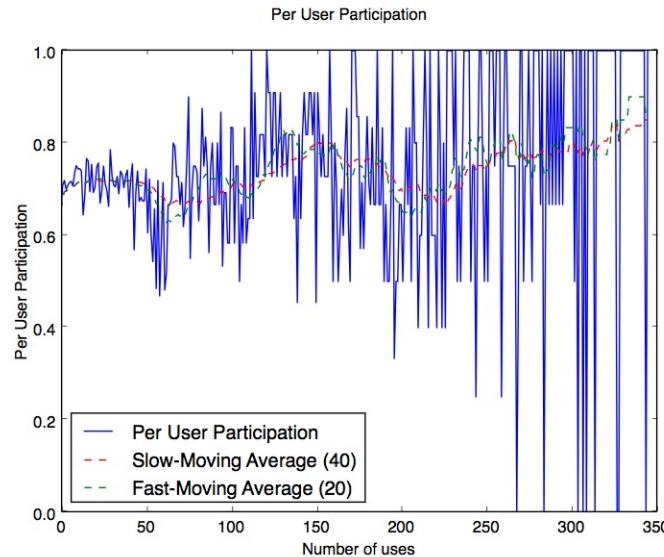


Figure 4.9: Participation Rate Per User

The two things that immediately stand out about these plots are the general upward drift and increasing volatility over time. The general upward drift of each conversation quality metric on a user-level supports the hypothesis that the UCB1-AKSB was increasing conversation quality over time. Another possible source of this

upward drift is that users who visit the site frequently (who I will refer to as power-users) are more interested in having high-quality conversations and will seek to do so independently of the conversation starter. However, even if the latter hypothesis is true, the UCB1-AKSB algorithm still succeeded in providing a set of initial conversation starters to these users to make them become power-users, so there is still evidence that the UCB1-AKSB algorithm was effective. Additionally, the steady increase in the conversation de-anonymization rate (Figure 4.7) suggests that the algorithm actually performed quite well on a per-user basis, despite the modest cumulative de-anonymization performance (Figure 4.3).

On the other hand, the increasing volatility over time of each conversation quality metric represents the stratification of users into different classes. An abrupt change from a participation rate of 1.0 to 0.0 in one use in any of the above per-user plots is most likely the result of a power-user being paired with an amateur-user where the amateur-user disconnects from the conversation before the power-user has a chance to participate meaningfully.

This user stratification is empirically validated by the behavior of TA users. For example, in Figure 4.10, most of the visit density cluster in the 1-2 visit category and the 26-200 visit category. This bimodal distribution of user behavior is even more prevalent in Figure 4.11, where the distribution is also clumped around short and long lengths of time spent on TA. This clearly supports the idea that, on both a amount-of-time and number-of-visits basis, TA users clumped into either power-users or amateur-users.

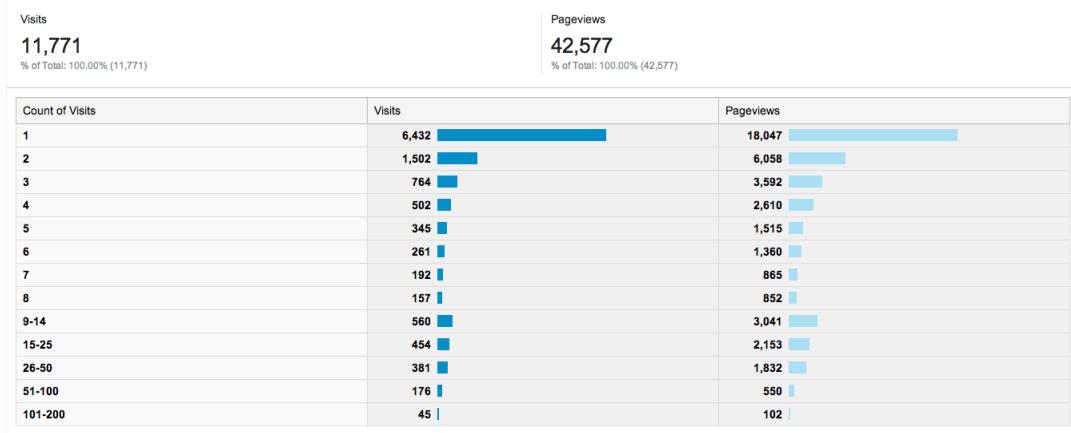


Figure 4.10: User Engagement by Visit Count

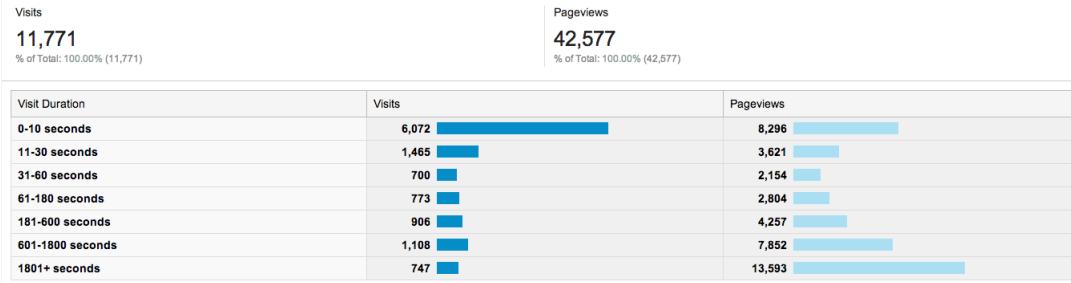


Figure 4.11: User Engagement by Visit Length

This user stratification into extremely high and extremely low involvement raises the issue of user saturation: people were either hooked and used the site extremely frequently or they used it a few times and left. As a result, the user base contained a small number of power-users and a large number of users who visited the site infrequently at best. This most likely resulted in a poor user experience for both classes of users: the amateur-users may not have become acclimated to the norms of anonymous conversation and might be intimidated by a conversation with a power-user, but bored with a conversation with another amateur-user. Conversely, a power-user would be used to the norms of an anonymous conversation, and would likely get bored conversing with amateur-user (who would still be operating under normal conversational norms, which the power-users are trying to escape). Although a power-

user would enjoy talking to another power-user, he or she would quickly run out of new power-users to talk to.

The solution to this problem is to have a ‘casual user’: one who can bridge the gap between the two previous user categories and provide a better TA experience for all. However, the problem that TA ran into with Princeton is the size of the potential user base: people at Princeton who would rather use TA over socializing with their current friends. This user base is not very large, but what if it were possible to combine such a user base from multiple schools into an even larger user base? This is the motivation for the project that I am currently working on: a larger version of Tigers Anonymous called Campus Anonymous, which will be released to all the universities in the Ivy League, thus broadening the user base and increasing the likelihood of attracting a non-trivial set of ‘casual users’.

Chapter 5

Conclusions

The purpose of this thesis is to test a new, weakly context-dependent multi-armed bandit algorithm that not only balances exploration and exploitation, but also selects never-before-seen arms for each user-context. By assuming that the reward from each arm is I.I.D, this new algorithm leverages all historical data (regardless of user-context) but still allows a certain level of user-specific recommendations. This, in turn, allows it to scale well over a large user base.

In order to test this new algorithm (UCB1-AKSB), I implemented it in Tigers Anonymous (TA) and optimized the algorithm to select conversation starters in order to maximize conversation de-anonymizations. After collecting and analyzing conversation and user data, we can draw the following conclusions about the UCB1-AKSB algorithm.

First, the data gathered suggests that TA violated some of the key assumptions made when applying the UCB1-AKSB algorithm. For example, the two different regimes in the cumulative regret analysis in Section 4.1 suggest that the distribution of payouts for each bandit arm were not I.I.D as assumed, but rather weakly Markovian. However, for the second regime, the algorithm behaved as expected and produced logarithmic cumulative regret (in line with the canonical UCB1 algorithm).

Moreover, the process of TA conversation de-anonymization may have been more context-dependent than UCB1-AKSB accounted for.

Second, it seems that UCB1-AKSB did not improve the overall conversation de-anonymization rate (although it improved the proportion of de-anonymizations given the option to de-anonymize). This is probably because conversation de-anonymization was not only dependent on the conversation starter and the binary variable of de-anonymization was not fine-grained enough to measure incremental improvements in the de-anonymization rate. However, UCB1-AKSB resulted in noticeable improvement in conversation quality over time on the basis of other metrics of conversation quality. In addition, UCB1-AKSB resulted in consistent improvement in per-user performance in all three of the conversation quality metrics used in this paper (conversation de-anonymization, total conversation length and participation rate).

Finally, the data suggests that, over time, the TA user base became stratified into two categories: power-users and amateur-users, with little middle ground. This resulted in increasing volatility in per-user conversation quality metrics (as explained in Section 4.3). It also resulted in user saturation, which decreased the overall TA user experience and likely contributed to the decline in the usage of the site.

In all, it seems that the UCB1-AKSB algorithm performed quite well given the violation of some of its fundamental assumptions, its calibration on an imperfect metric and the stratification/saturation of its user base. However, this thesis was only a single empirical test of the new algorithm, which means that there is substantial room for further testing of this algorithm. For future research in this area, this thesis can make some important recommendations.

First, any future research would benefit from a larger user base to avoid the problems of user saturation and stratification mentioned in Section 4.3. This would make the formation of a ‘casual-user’ class more likely, which would bridge the gap between amateur and power-users, thus improving the user experience and making

the chatroom gain more traction with the user base.

Second, the algorithm’s performance could be improved by using a better measure of conversation quality than a simple binary variable (i.e. the likelihood of de-anonymization). For example, one could create a hybrid metric by assigning a more finely tuned score based on multiple metrics (i.e. 0 for immediate disconnect, 0.3 for a short conversation, 0.7 for a long conversation, and 1 for a de-anonymization). Another option would remove the binary variable (i.e. de-anonymization) altogether, and instead use another metric (such as the participation rate or average conversation length).

Finally, the algorithm itself could be modified to work more effectively with a static set of arms (or a dynamic set of arms with a slow turnover rate). The current version of the algorithm outlined in Chapter 3 simply serves a random arm if both user pairs have already seen the conversation starter. This could be improved by implementing a sliding time window for context-dependency: that is, only use the conversation starters that users u and v haven’t seen within the last week or month. This could also be implemented as a sliding number of plays (i.e. only pick from arms that users u and v haven’t seen within the last 10 plays).

Not only is there room for future research on the UCB1-AKSB algorithm, but it has broad applications beyond simply selecting quirky conversation starters for an anonymous chatroom. For starters, it is very fast and computationally easy to implement, as opposed to more intensive contextual bandit algorithms such as LinUCB in Li et al. (2010). In addition, UCB1-AKSB (or at least some of its principles) can be applied to nearly every bandit recommendation algorithm. In stochastic bandit recommendation systems that would benefit from a degree of context-dependency (i.e. making novel recommendations to a homogenous user base), the UCB1-AKSB would provide just enough context dependency while still taking advantage of the assumption that user responses are I.I.D. It is easy to see this algorithm being applied

to recommend news articles or movies of a certain genre, where novelty is important but users are relatively homogenous. In other bandit algorithms that are more heavily context dependent, the arm-filtering mechanism used by UCB1-AKSB could be added on top of the existing bandit mechanism to ensure that users don't see an arm more than once.

In conclusion, the UCB1-AKSB algorithm fills a gap in the literature between stochastic multi-armed bandits and context-dependent bandits. It provides a solution to a real-world problem in Tigers Anonymous and could be applied in a variety of other contexts. By providing a weakly context-dependent multi-armed bandit algorithm and showing its potential for performance in the real-world, I hope that this thesis will help spur future research into this new class of algorithms.

Appendix A

Data Analysis Code

The following is the code used to process the TA conversation data and produce the figures in this thesis.

```
import json
from datetime import datetime, date, timedelta
import collections
import matplotlib.pyplot as plt
import matplotlib.dates as mdates

#####
### Helper functions to process input JSON file ###
#####

def conversation_date(conv):
    return date.fromtimestamp(conv['startTime']/1000.0)

def first_user(conv):
    return conv['userID1']

def second_user(conv):
    return conv['userID2']

def users(conv):
    return [first_user(conv), second_user(conv)]

def first_user_clicked(conv):
    return conv['user1Clicked']

def second_user_clicked(conv):
    return conv['user2Clicked']

def fb_match_occurred(conv):
    return first_user_clicked(conv) and second_user_clicked(conv)

def first_user_messages_sent(conv):
    return conv['user1MessagesSent']

def second_user_messages_sent(conv):
```

```

        return convo['user2MessagesSent']
def total_user_messages_sent(convos):
    return (first_user_messages_sent(convos) +
            second_user_messages_sent(convos))
def first_user_participated(convos):
    return first_user_messages_sent(convos) != 0
def second_user_participated(convos):
    return second_user_messages_sent(convos) != 0
def no_immediate_disconnect_occurred(convos):
    return (first_user_participated(convos) or
            second_user_participated(convos))

#####
### Helper functions to extract data from input JSON file #####
#####

# Returns list of tuples of form (date, {wins: 0, plays: 0})
def daily_win_play_data(data, win_metric):
    regret_data = {}
    for convo in data:
        date = conversation_date(convos)
        if date in regret_data:
            regret_data[date]['wins'] += win_metric(convos)
            regret_data[date]['plays'] += 1.0
        else:
            regret_data[date] = {'wins': 0.0 + win_metric(convos), 'plays': 1.0}
    return sorted(regret_data.items())

# Returns list of tuples of form (date, {wins: 0, plays: 0})
def cumulative_win_play_data(data, win_metric):
    win_play_data = daily_win_play_data(data, win_metric)
    output = {}
    output[win_play_data[0][0]] = win_play_data[0][1]
    for i in range(1, len(win_play_data)):
        curr_data_point = win_play_data[i]
        curr_date = curr_data_point[0]
        curr_wins = curr_data_point[1]['wins']
        curr_plays = curr_data_point[1]['plays']

        prev_date = win_play_data[i-1][0]
        prev_wins = output[prev_date]['wins']
        prev_plays = output[prev_date]['plays']

        output[curr_date] = {'wins': curr_wins + prev_wins, 'plays':
            curr_plays + prev_plays}
    return sorted(output.items())

```

```

# Returns list of cumulative regret values with respect to the win_metric
def cumulative_regret(data, win_metric):
    def reduce_function(total, play):
        try:
            curr_wins = total[-1][1]
            curr_plays = total[-1][2]
        except:
            curr_wins = 0
            curr_plays = 0
        total.append((play, curr_wins + play, curr_plays + 1))
    return total
    results = reduce(reduce_function, map(lambda convo: win_metric(convo),
                                           data), [])
    optimal_win_ratio = results[-1][1]/float(results[-1][2])
    return map(lambda (result, wins, plays): (optimal_win_ratio * plays) -
               wins, results)

# Returns ordered dictionary of form {date: win_ratio}
def win_ratio(win_play_data, threshold=0):
    output = {}
    for data_point in win_play_data:
        date = data_point[0]
        wins = data_point[1]['wins']
        plays = data_point[1]['plays']
        if plays >= threshold:
            output[date] = wins/plays
    return collections.OrderedDict(sorted(output.items()))

# Generates a list where the (i-1)-th index represents
# the average value of the user1/2_metric for all the users
# on their i-th play conditional on having played i or more times.
# For example, a result of 0.5 in the list index of 9 would mean
# that for all users who played 10 or more times, the average value
# of the metric on their 10th play was 0.5.
def generate_user_data(data, user1_metric, user2_metric):
    user_data = []
    for convo in data:
        # Because migration to new server caused unique userIDs to
        # reset after 1/29/2014
        if conversation_date(convo) > date(2014, 1, 29):
            if first_user(convo) in user_data:
                user_data[first_user(convo)].append(user1_metric(convo))
            else:
                user_data[first_user(convo)] = [user1_metric(convo)]
            if second_user(convo) in user_data:
                user_data[second_user(convo)].append(user2_metric(convo))
            else:

```

```

        user_data[second_user(convos)] = [user2_metric(convos)]
user_aggregate = []
for user in user_data:
    user_aggregate.append(user_data[user])
user_aggregate = map(None, *user_aggregate)
output = []
for iteration in user_aggregate:
    num_successes = reduce(lambda a, b: a + b if b != None else a,
                           iteration, 0.0)
    num_trials = reduce(lambda a, b: a + 1 if b != None else a, iteration,
                        0)
    output.append(num_successes/num_trials)
return output

# Given data in the form of a list of values, return a list containing
# the simple moving average with a specified window_size
def get_moving_average(data, window_size):
    output = [data[0]]
    for i in range(1, window_size-1):
        output.append(sum(data[0:i+1])/float(i+1))
    for i in range(window_size, len(data)+1):
        output.append(sum(data[i-window_size:i])/float(window_size))
    return output

#####
### Helper functions to format and display the data ###
#####

# Takes in ordered dictionary d and saves it as name
def save_dict_plot(d, title, xlabel='Date', ylabel='Metric'):
    x = []
    y = []
    for key in d:
        x.append(key)
        y.append(d[key])
    fig, ax = plt.subplots()
    fig.suptitle(title)
    plt.xlabel(xlabel)
    plt.ylabel(ylabel)
    plt.plot(x, y)
    fig.autofmt_xdate()
    fig.savefig('../Figures/' + title.replace(' ', '_') + '.jpg')

def save_list_plot(l, title, xlabel='Conversation Index', ylabel='Metric'):
    fig, ax = plt.subplots()
    fig.suptitle(title)
    plt.xlabel(xlabel)

```

```

plt.ylabel(ylabel)
plt.plot(l)
fig.savefig('..../Figures/' + title.replace(' ', '_') + '.jpg')

# Win ratio analysis, both daily and cumulatively
def do_win_ratio_analysis(data, metric, metric_label, y_label,
    threshold=20):
    daily_wins_and_plays = daily_win_play_data(data, metric)
    cumulative_wins_and_plays = cumulative_win_play_data(data, metric)
    daily_win_ratio = win_ratio(daily_wins_and_plays, threshold)
    cumulative_win_ratio = win_ratio(cumulative_wins_and_plays)
    save_dict_plot(daily_win_ratio, 'Daily ' + metric_label, ylabel=y_label)
    save_dict_plot(cumulative_win_ratio, 'Cumulative ' + metric_label,
        ylabel=y_label)

def do_cumulative_regret_analysis(data, metric, metric_label,
    y_label='Cumulative Regret'):
    cumulative_regret_data = cumulative_regret(data, metric)
    save_list_plot(cumulative_regret_data, metric_label, ylabel=y_label)

def do_per_user_analysis(data, user1_metric, user2_metric, metric_title,
    slow_moving_average_window=40, fast_moving_average_window=20,
    legend_location="upper left"):
    user_data = generate_user_data(data, user1_metric, user2_metric)
    slow_moving_average = get_moving_average(user_data,
        slow_moving_average_window)
    fast_moving_average = get_moving_average(user_data,
        fast_moving_average_window)
    fig, ax = plt.subplots()
    fig.suptitle(metric_title)
    plt.xlabel('Number of uses')
    plt.ylabel(metric_title)
    plt.plot(user_data, label=metric_title)
    plt.plot(slow_moving_average, 'r--', label=('Slow-Moving Average (' +
        str(slow_moving_average_window) + ')'))
    plt.plot(fast_moving_average, 'g--', label=('Fast-Moving Average (' +
        str(fast_moving_average_window) + ')'))
    plt.legend(loc=legend_location)
    fig.savefig('..../Figures/' + metric_title.replace(' ', '_') + '.jpg')

#####
### Commands to use the above functions to perform data analysis #####
#####

# Load the JSON data
data = json.load(open('ta_data.json', 'r'))

```

```

# Do win-ratio analysis
do_win_ratio_analysis(data, fb_match_occurred, 'TA De-Anonymization',
                      'Proportion of De-Anonymized Conversations')
do_win_ratio_analysis(data, no_immediate_disconnect_occurred, 'TA
                      Participation', 'Participation Rate')
do_win_ratio_analysis(data, total_user_messages_sent, 'TA Conversation
                      Length (Messages Exchanged)', 'Average Conversation Length (# of
                      messages exchanged)')

# Do cumulative regret analysis
do_cumulative_regret_analysis(data, fb_match_occurred, 'TA
                               De-Anonymization Regret')
do_cumulative_regret_analysis(data, no_immediate_disconnect_occurred, 'TA
                               Participation Rate Regret')
do_cumulative_regret_analysis(data, total_user_messages_sent, 'TA Messages
                               Sent Regret')

# Do per-user analysis
do_per_user_analysis(data, first_user_clicked, second_user_clicked, 'Per
                      User De-Anonymization')
do_per_user_analysis(data, first_user_participated,
                     second_user_participated, 'Per User Participation',
                     legend_location="lower left")
do_per_user_analysis(data, first_user_messages_sent,
                     second_user_messages_sent, 'Per User Messages Sent')

```

Appendix B

Conversation Starters

The following is the complete set of conversation starters used in Tigers Anonymous. I personally edited and approved each question on this list specifically for Tigers Anonymous, and my inspiration for questions on this list came from my own Princeton experience, conversations with friends and a variety of public-domain online sources.

```
exports.list = [
    "What animal is your Patronus?",
    "If you ruled the world, what laws would you make?",
    "What was your last dream about?",
    "What would you do if you won the lottery?",
    "What does your dream house look like?",
    "What was your favorite vacation?",
    "If you could go back in time to change one thing what would it be?",
    "What's the greatest invention of all time?",
    "Have you ever been admitted to the hospital?",
    "Have you ever had any brushes with the law?",
    "What's the best practical joke you've played on someone?",
    "What's the best practical joke someone's pulled on you?",
    "What is your best achievement to date?",
    "If you could live anywhere, where would it be?",
    "What's your favorite song?",
    "What's your favorite word (inappropriate or otherwise)?",
    "What's the longest period of time you've gone without sleep?",
    "Do you have any scars?",
    "If you could change anything about yourself what would it be?",
    "Would you rather trade some intelligence for looks or looks for
```

intelligence?",
"Have you ever had a secret admirer?",
"If you could ask your future self one question, what would it be?",
"Are you a good liar?",
"What's your favorite joke?",
"What's the worst present you've ever gotten?",
"What's your favourite saying?",
"Have you ever accidentally injured anyone?",
"What cartoon character would you love to see in 21st century life?",
"What's the word you use most often?",
"What's your dream job?",
"Which song annoys you the most?",
"What's your first thought when you wake up?",
"If you could steal one thing in the world, what would it be?",
"What's your favorite Pokemon character?",
"When did you stop believing in Santa?",
"What's your favorite Disney movie?",
"What's your life motto?",
"What's the most unusual thing you've ever eaten?",
"Do you collect anything?",
"What thing would you like to bring back into fashion?",
"What makes you nervous?",
"What's the worst pickup line you've ever heard?",
"What do you do when you forget someone's name immediately after they've introduced themselves?",
"Have you ever been in a fight?",
"Have you ever started a rumor?",
"What's the most memorable rumor you've heard about yourself?",
"Is there anything about the opposite sex you just don't understand?",
"If you had a warning label, what would yours say?",
"Which fictional character do you wish was real?",
"Who was your first crush?",
"Do you believe in destiny or free will?",
"What's the best decision you've made so far?",
"Who would you want to be with on a desert island?",
"If you had to pick a new name for a week, what would it be?",
"What is your first memory?",
"Where did you go on your first ride on an airplane?",
"Who was your first best friend?",
"What was your first detention for?",
"What would be the name of your debut solo album?",
"What's something you get compulsive about?",
"Have you ever stolen anything?",
"What was the last social faux pas you made?",
"What makes you nostalgic?",
"What's the scariest thing you've ever done?",
"What fairy tale character do you most associate with?",

"What's the craziest thing you've ever done for someone?",
"What's the best piece of advice anyone has ever given you?",
"Do you have a Princeton bucket list?",
"What's your favorite memory at Princeton?",
"What building would you donate to Princeton?",
"What is one thing you always wanted as a kid, but never got?",
"What is the nicest thing someone else has done for you?",
"If you could time travel, what would you do?",
"If you went to a psychiatrist, what would he/she say you suffer from?",
"What one thing annoys you most at a restaurant?",
"What do Princeton students do too much of today?",
"What would you like to spend more time doing?",
"If you could dis-invent one thing, what would it be?",
"How would you dispose of a dead body?",
"What's the most recent dream you can remember?",
"What's something about you that people wouldn't expect?",
"If you could change one thing about the world, what would it be?",
"What's your favorite genre of music?",
"If you could eat lunch with one famous person, who would it be?",
"How are you feeling right now?",
"What do you think about the most?",
"Do you sing in the shower?",
"Before Princeton, what did you want to be when you grew up?",
"What is your best childhood memory?",
"What's something embarrassing that happened to you?",
"If you could live in any city in the world, where would it be?",
"Where do you want to travel to?",
"What's something spontaneous that you've done?",
"If you could only eat one food for the rest of your life, what would it be?",
"What's your biggest pet peeve?",
"What was the happiest moment in your life?",
"What quality about yourself do you value most?",
"What are you most proud about in your life?",
"What is your biggest concern about the future?",
"What is the biggest lesson you've learned in life thus far?",
"Do you think people can control their own destinies?",
"How is your relationship with your parents?",
"If you could go back and relive a day in your life, what would you change?",
"What is the weirdest thing about you?",
"If you could have any superpower, which one would you pick?",
"What is the last thing you do before you go to sleep?",
"Whats the first thing you notice when you meet someone new?",
"Whats one of your worst habits?",
"If your house was on fire, what's the one thing you'd want to take with you?",

"If money was no object, what would you be doing with your life?",
"What does your vision of a utopian society look like?",
"If you only had one day left to live, what would you do?",
"What's one thing that you learned this week?",
"What was the last thing you thought about last night?",
"What were you like as a kid?",
"What is one thing you miss about being a kid?",
"Do you believe in soul mates?",
"Do you believe in love at first sight?",
"What's one thing you'd like to change about Princeton?",
"How was your RCA during your freshman year?"

];

Appendix C

TA Back-End Implementation

The following pieces of code implement the back-end and front-end functionality of Tigers Anonymous unrelated to the UCB1-AKSB algorithm.

C.1 Princeton IP-Address Filtering Functionality

```
var range_check = require('range_check');

// Pre-defined Princeton IP address blocks
var princetonIPs = [
    "128.112.0.0/16",
    "140.180.0.0/16",
    "204.153.48.0/22",
    "66.180.176.0/24",
    "66.180.177.0/24",
    "66.180.180.0/22"
];

// Check to ensure that the user's IP is a valid Princeton IP
var isValidIP = function (userIP) {
    if (userIP === "127.0.0.1" || // for debugging
        range_check.in_range(userIP, "192.168.0.0/16") ||
        range_check.in_range(userIP, "10.0.0.0/8")) {
        return true;
    }
    for (var i = 0; i < princetonIPs.length; i++) {
        if (range_check.in_range(userIP, princetonIPs[i])) {
```

```

        return true;
    }
}

return false;
}

exports.isValidIP = isValidIP;

```

C.2 User Matching Functionality

```

var mongoose = require('mongoose');
var Conversation = mongoose.model('Conversation');
var ucb = require('./ucb');
var mailer = require('./mailer');

function User(socket, userID) {
    this.socket = socket;
    this.id = userID;
    this.partner = null;
    this.conversation = null;
    this.buttonClicked = false;
    this.messagesSent = 0;
    this.name = null;
    this.fbLink = null;

    var user = this;
    this.socket.on('disconnect', function() {
        if (!user.conversation) return;

        if (!user.conversation.endTime) {
            user.conversation.chatLog.push({
                date: new Date(),
                user: '',
                text: '*** ' + user.pseudonym + ' disconnected ***'
            });
        }

        user.conversation.endTime = new Date();
        user.conversation.save();
        user.partner.socket.emit('finished');
        user.partner.socket.disconnect();
    })
}

```

```

this.socket.on('chat message', function(data) {
  if (!user.conversation) return;

  user.conversation.chatLog.push({
    date: new Date(),
    user: user.pseudonym,
    text: data.message
  });

  user.messagesSent++;
  user.socket.emit('chat message', {
    name: 'You',
    message: data.message
  });
}

var userName = user.conversation.revealed ? user.name : 'Anonymous
Tiger';
user.partner.socket.emit('chat message', {
  name: userName,
  message: data.message
});
});

this.socket.on('dropdown displayed', function(data) {
  if (!user.conversation) return;

  user.conversation.buttonDisplayed = true;
});

this.socket.on('identity', function(data) {
  if (!user.conversation) return;

  user.conversation.chatLog.push({
    date: new Date(),
    user: '',
    text: '*** ' + user.pseudonym + ' accepted dropdown ***'
  });

  user.name = data.name;
  user.fbLink = data.link;
  user.buttonClicked = true;

  if (user.partner.buttonClicked) {
    user.socket.emit('reveal', {
      name: user.partner.name,
      link: user.partner.fbLink
    });
  }
});

```

```

        user.partner.socket.emit('reveal', {
            name: user.name,
            link: user.fbLink
        });
        user.conversation.revealed = true;

        user.conversation.chatLog.push({
            date: new Date(),
            user: '',
            text: '*** Facebook identities revealed ***'
        });
    }
});

this.socket.on('typing', function() {
    if (!user.conversation) return;

    user.partner.socket.emit('typing');
});

this.socket.on('not typing', function() {
    if (!user.conversation) return;

    user.partner.socket.emit('not typing');
});

function ConversationWrapper() {
    this.user1 = null;
    this.user2 = null;
    this.startTime = new Date();
    this.endTime = null;
    this.question = null;
    this.buttonDisplayed = false;
    this.revealed = false;
    this.chatLog = [];

    var self = this;
    this.save = function() {
        new Conversation({
            userID1: self.user1.id,
            userID2: self.user2.id,
            question: self.question,
            startTime: self.startTime,
            endTime: self.endTime,
            buttonDisplayed: self.buttonDisplayed,
            user1Clicked: self.user1.buttonClicked,

```

```

        user2Clicked: self.user2.buttonClicked,
        user1MessagesSent: self.user1.messagesSent,
        user2MessagesSent: self.user2.messagesSent
    }).save();

    if (process.env.NODE_ENV === 'production') {
        mailer.sendMail(this);
    }
};

var queue = new Array();
exports.connectChatter = function(socket, userID) {
    var user = new User(socket, userID);

    user.socket.emit('entrance');
    user.socket.emit('waiting');

    if (queue.length === 0) {
        queue.push(user);

        // TODO: remove listener instead of checking index
        user.socket.on('disconnect', function() {
            var index = queue.indexOf(user);
            if (index !== -1) {
                queue.splice(index, 1);
            }
        });
    } else {
        var conversation = new ConversationWrapper();
        conversation.user1 = user;
        user.conversation = conversation;
        user.pseudonym = 'Origin';

        var partner = queue.shift();
        user.partner = partner;
        partner.partner = user;
        conversation.user2 = partner;
        partner.conversation = conversation;
        partner.pseudonym = 'Black';

        ucb.getQuestion(Conversation, user, partner, function(question) {
            user.conversation.question = question;
            user.socket.emit('matched', {
                question: question
            });
            partner.socket.emit('matched', {

```

```

        question: question
    });

    conversation.chatLog.push({
        date: new Date(),
        user: '',
        text: question
    });
}
};


```

C.3 Web Server Functionality

```

var express = require('express'),
app = express(),
server = require('http').createServer(app),
io = require('socket.io').listen(server);
mongoose = require('mongoose'),
princeton = require('./server/princeton'),
conversation = require('./server/conversation'),
chatter = require('./server/chatter');

var port = process.env.PORT || 5000;
server.listen(port);

var mongoUrl;
io.configure('development', function() {
    mongoUrl = 'mongodb://localhost/test';
});
io.configure('production', function() {
    mongoUrl = process.env.MONGOHQ_URL;
});
mongoose.connect(mongoUrl);

var connectedUsers = {};

app.get('/count', function(req, res) {
    var count = Object.keys(connectedUsers).length;
    res.send(count.toString());
});

io.configure('production', function() {

```

```

io.set('log level', 1);
io.set('transports', ['websocket']);

io.set('authorization', function(handshakeData, callback) {
    // Check if Princeton IP
    var ipAddr = getClientIP(handshakeData);
    var isValidIP = princeton.isValidIP(ipAddr);
    if (!isValidIP) {
        callback('Sorry, this site is only for Princeton students!', false);
        return;
    }

    // Check if already connected to server
    if (ipAddr in connectedUsers) {
        callback('Sorry, you can only chat with one person at a time!', false);
        return;
    }

    callback(null, true);
});

// Needed to get the client's IP on Heroku for socket.io
function getClientIP(handshakeData) {
    var forwardedIps = handshakeData.headers['x-forwarded-for'];
    if (forwardedIps) {
        return forwardedIps.split(', ')[0];
    } else {
        return handshakeData.address.address;
    }
}

function getValueFromCookie(name, cookie) {
    if (cookie) {
        var pairs = cookie.split('; ');
        for (var i = 0; i < pairs.length; i++) {
            var pair = pairs[i].split('=');
            if (pair[0] === name) {
                return pair[1];
            }
        }
    }
}

io.sockets.on('connection', function(socket) {

```

```
var userID = getValueFromCookie('userID',
    socket.handshake.headers.cookie);
if (userID) {
    // Add user to list of connected users
    var ipAddr = getClientIP(socket.handshake);
    connectedUsers[ipAddr] = true;
    socket.on('disconnect', function() {
        delete connectedUsers[ipAddr];
    });
}

chatter.connectChatter(socket, userID);
} else {
    socket.disconnect();
}
});
```

C.4 Conversation Metadata Logging Model

```
var mongoose = require('mongoose');

var conversationSchema = new mongoose.Schema({
    userID1: String,
    userID2: String,
    startTime: Date,
    endTime: Date,
    question: String,
    buttonDisplayed: Boolean,
    user1Clicked: Boolean,
    user2Clicked: Boolean,
    user1MessagesSent: Number,
    user2MessagesSent: Number
});

mongoose.model('Conversation', conversationSchema);
```

Appendix D

TA Front-End Implementation

D.1 Homepage

The homepage (shown below in D.1) is implemented with the code shown at the bottom of this section.

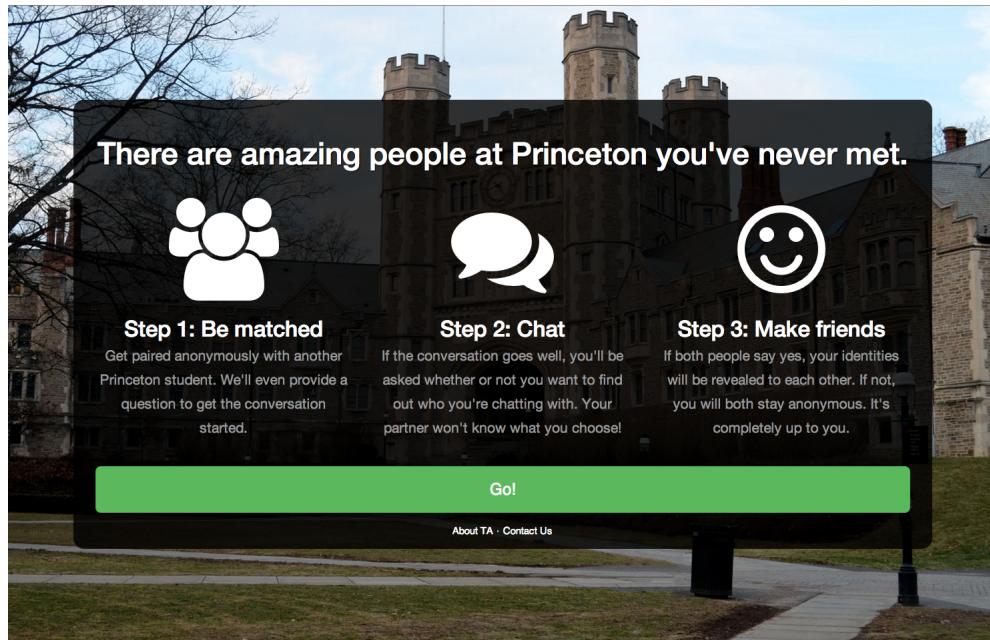


Figure D.1: Tigers Anonymous Homepage

```

<html>
  <head prefix="og: http://ogp.me/ns#">
    <title>Tigers Anonymous</title>
    <link rel="icon" href="img/favicon.ico" type="image/x-icon">
    <meta name="viewport" content="width=device-width, initial-scale=1.0,
      user-scalable=no">
    <meta property="og:title" content="Tigers Anonymous">
    <meta property="og:description" content="There are amazing people at
      Princeton you've never met.">
    <meta property="og:url" content="http://www.tigersanonymous.com">
    <meta property="og:image"
      content="http://www.tigersanonymous.com/img/ta1024.png">
    <meta property="og:image:type" content="image/png">
    <meta property="og:image:width" content="1024">
    <meta property="og:image:height" content="1024">
    <link rel="stylesheet"
      href="//netdna.bootstrapcdn.com/bootstrap/3.0.3/css/bootstrap.min.css">
    <link
      href="//netdna.bootstrapcdn.com/font-awesome/4.0.3/css/font-awesome.css"
      rel="stylesheet">
    <link href="css/index.css" rel="stylesheet" type="text/css"
      media="all">
    <script>
      (function(i,s,o,g,r,a,m){i['GoogleAnalyticsObject']=r;i[r]=i[r]||function(){
        (i[r].q=i[r].q||[]).push(arguments)},i[r].l=1*new
        Date();a=s.createElement(o),
        m=s.getElementsByTagName(o)[0];a.async=1;a.src=g;m.parentNode.insertBefore(a,m)
      })(window,document,'script','//www.google-analytics.com/analytics.js','ga');
      ga('create', 'UA-23357698-2', 'tigersanonymous.com');
      ga('send', 'pageview');
    </script>
  </head>
  <body class="cover">
    <div class="wrapper">
      <div class="container">
        <div class="row text-center">
          <div class="col-md-12">
            <h1 class="hook">There are amazing people at Princeton you've
              never met.</h1>
          </div>
        </div>
        <div class="row how-it-works">
          <div class="col-md-4">
            <div class="row text-center padded-icon">
              <i class="fa fa-users large-icon"></i>
            </div>
            <div class="row text-center padded-text">

```

```

<h2>
  Step 1: Be matched<br>
  <small>Get paired anonymously with another Princeton
  student. We'll even provide a question to get the
  conversation started.</small>
</h2>
</div>
</div>
<div class="col-md-4">
  <div class="row text-center padded-icon">
    <i class="fa fa-comments large-icon"></i>
  </div>
  <div class="row text-center padded-text">
    <h2>
      Step 2: Chat<br>
      <small>If the conversation goes well, you'll be asked
      whether or not you want to find out who you're chatting
      with. Your partner won't know what you choose!</small>
    </h2>
    </div>
  </div>
  <div class="col-md-4">
    <div class="row text-center padded-icon">
      <i class="fa fa-smile-o large-icon"></i>
    </div>
    <div class="row text-center padded-text">
      <h2>
        Step 3: Make friends<br>
        <small>If both people say yes, your identities will be
        revealed to each other. If not, you will both stay
        anonymous. It's completely up to you.</small>
      </h2>
      </div>
    </div>
  </div>
  <div class="row">
    <div class="col-md-12">
      <a href="/chat" class="go-btn btn btn-success btn-xlg
        btn-block">Go!</a>
    </div>
  </div>
  <div class="row text-center">
    <div class="col-md-12 footer">
      <a href="/about">About TA </a>
      &#8901;
      <a href="mailto:originblack609@gmail.com"> Contact Us</a>
    </div>
  </div>

```

```
</div>
</div>
</div>
</body>
</html>
```

D.2 About Page

The About page (shown below in D.2) is implemented with the code shown at the bottom of this section.

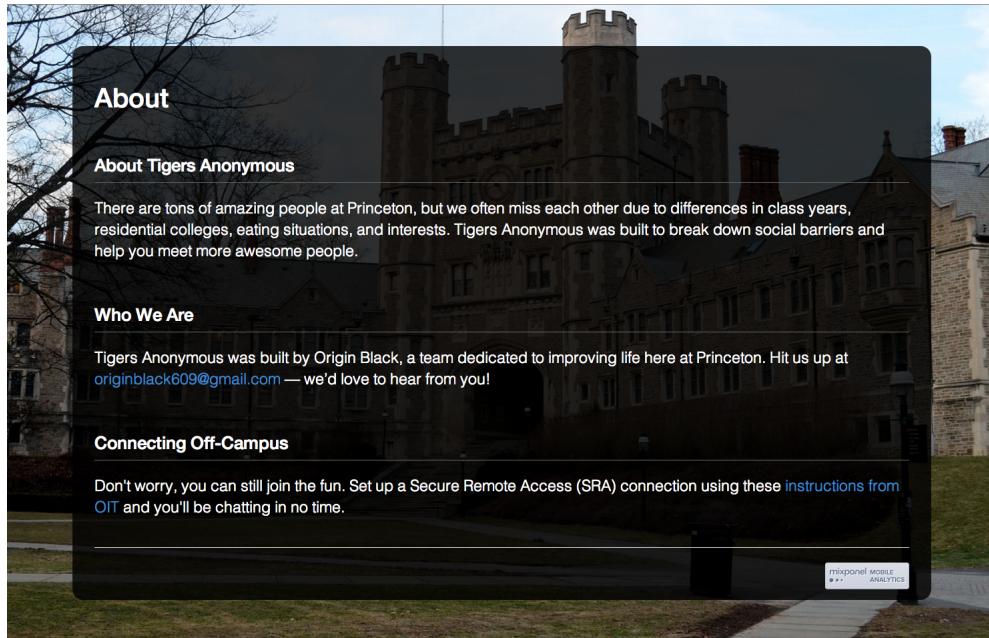


Figure D.2: Tigers Anonymous About Page

```
<!DOCTYPE html>
<html>
  <head>
    <title>About - Tigers Anonymous</title>
    <link rel="icon" href="img/favicon.ico" type="image/x-icon">
    <meta name="viewport" content="width=device-width, initial-scale=1.0,
      user-scalable=no">
    <link rel="stylesheet"
      href="//netdna.bootstrapcdn.com/bootstrap/3.0.3/css/bootstrap.min.css">
```

```

<link href="css/index.css" rel="stylesheet" type="text/css"
      media="all">
<script>
(function(i,s,o,g,r,a,m){i['GoogleAnalyticsObject']=r;i[r]=i[r]||function(){
(i[r].q=i[r].q||[]).push(arguments)},i[r].l=1*new
Date();a=s.createElement(o),
m=s.getElementsByTagName(o)[0];a.async=1;a.src=g;m.parentNode.insertBefore(a,m)
})(window,document,'script','//www.google-analytics.com/analytics.js','ga');
ga('create', 'UA-23357698-2', 'tigersanonymous.com');
ga('send', 'pageview');
</script>
</head>
<body class="cover">
<div class="wrapper">
<div class="container">
<div class="row">
<div class="col-md-12">
<h1>About
<div class="header" id="about">
<h3>About Tigers Anonymous</h3>
</div>
<p class="lead">
There are tons of amazing people at Princeton, but we often
miss each other due to differences in class years,
residential colleges, eating situations, and interests.
Tigers Anonymous was built to break down social barriers and
help you meet more awesome people.
</p>
<div class="header" id="whowearare">
<h3>Who We Are</h3>
</div>
<p class="lead">
Tigers Anonymous was built by Origin Black, a team dedicated to
improving life here at Princeton. Hit us up at <a
href="mailto:originblack609@gmail.com">originblack609@gmail.com</a>
&#8212; we'd love to hear from you!
</p>
<div class="header" id="offcampus">
<h3>Connecting Off-Campus</h3>
</div>
<p class="lead">
Don't worry, you can still join the fun. Set up a Secure Remote
Access (SRA) connection using these <a
href="http://helpdesk.princeton.edu/kb/display.plx?ID=6023">instructions
from OIT</a> and you'll be chatting in no time.
</p>
</div>

```

```
</div>
<div class="row">
  <div class="col-md-12 text-right">
    <hr>
    <a href="https://mixpanel.com/f/partner"></a>
  </div>
</div>
</div>
</body>
</html>
```

D.3 Chatroom

The Tigers Anonymous chatroom (shown below in D.3) is implemented with the code shown at the bottom of this section.

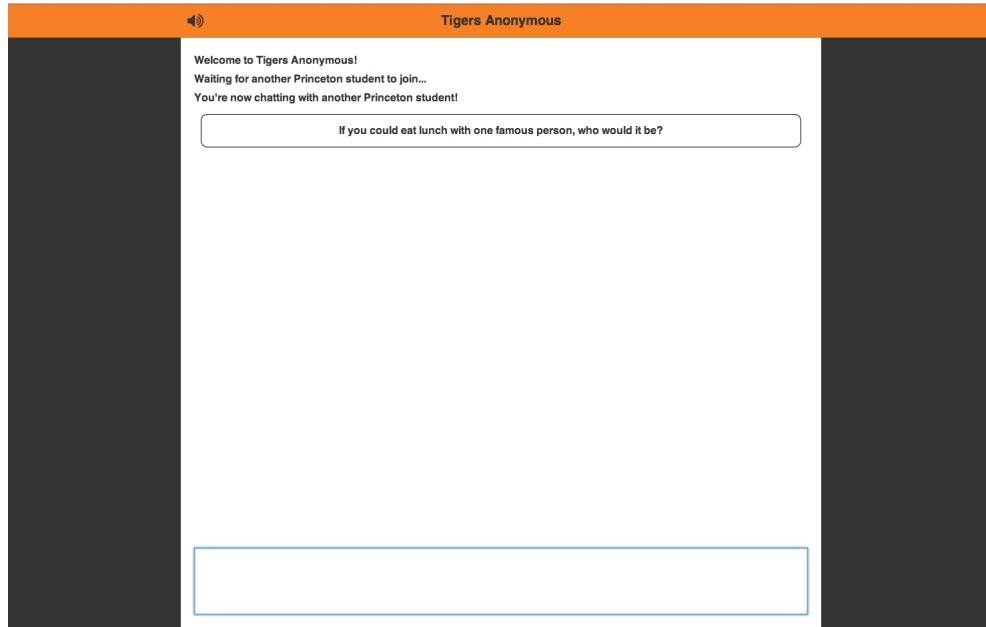


Figure D.3: Tigers Anonymous Chatroom

```
<!DOCTYPE html>
<html ng-app="pom">
  <head ng-controller="TitleCtrl">
```

```

<title ng-bind="getTitle()">Tigers Anonymous</title>
<link rel="icon" href="img/favicon.ico" type="image/x-icon">
<meta name="viewport" content="width=device-width, initial-scale=1.0,
    user-scalable=no">
<meta name="apple-mobile-web-app-capable" content="yes">
<link href="//netdna.bootstrapcdn.com/font-awesome/4.0.3/css/font-awesome.css"
    rel="stylesheet">
<link href="css/chat.css" rel="stylesheet" type="text/css" media="all">
</head>
<body ng-controller="ChatCtrl">
<div id="fb-root"></div>
<div class="nav">
<div class="nav-container">
    <span class="brand" href="/">Tigers Anonymous</span>
    <a class="volume" ng-click="playSound = !playSound" ng-cloak>
        <i class="fa fa-volume-up" ng-show="playSound"></i>
        <i class="fa fa-volume-off" ng-show="!playSound"></i>
    </a>
    <a class="circle-down" ng-show="dropdown.shouldShowMinimized() &&
        state == 'chatting'" ng-click="dropdown.show()" ng-cloak>
        <i class="fa fa-chevron-circle-down"></i>
    </a>
</div>
</div>
<div class="chat-container">
    <div class="dropdown" ng-show="dropdown.shouldShowFull() && state ==
        'chatting'" ng-cloak>
        <div class="question">
            Do you want to find out who you've been chatting with?<br>
            <span class="promise">We'll never post to Facebook without your
                permission. Promise.</span>
        </div>
        <div class="options">
            <button type="button" class="yes-btn"
                ng-click="dropdown.accept()">Yes</button>
            <button type="button" class="hide-btn"
                ng-click="dropdown.hide()">Hide</button>
        </div>
    </div>
    <div class="chatroom" pom-scroll-glue>
        <ul ng-cloak>
            <li ng-repeat="message in messages" ng-class="message.type"
                ng-switch="message.type">
                <div ng-switch-when="chat">
                    <span ng-class="{userName: !message.isPartner, partnerName:
                        message.isPartner}">{{message.name}}:</span>

```

```

<span ng-bind-html="message.text | linky |
    linkyNewlines"></span>
</div>
<div ng-switch-when="system">
<div ng-switch="message.template" ng-class="{important:
    message.important}">
<div ng-switch-when="entrance">
    Welcome to Tigers Anonymous!
</div>
<div ng-switch-when="waiting">
    Waiting for another Princeton student to join...
</div>
<div ng-switch-when="matched">
    You're now chatting with another Princeton student!<br>
    <div class="question-box">
        {{message.question}}
    </div>
</div>
<div ng-switch-when="selfRevealed">
    Your partner's identity will be revealed if they also want
    to discover yours.
</div>
<div ng-switch-when="partnerRevealed">
    Congratulations! You get to find out your partner's
    identity!<br>
    You've been chatting with: <a
        href="{{message.partnerLink}}"
        target="_blank">{{message.partnerName}}</a>
</div>
<div ng-switch-when="fbError">
    Sorry, there was an error connecting to Facebook. Please
    try again.
</div>
<div ng-switch-when="fbFake">
    Sorry, it looks like you're using a fake Facebook account.
</div>
<div ng-switch-when="finished">
    {{partnerName}} has disconnected. Refresh the page to
    start another chat!<br>
    What do you think about Tigers Anonymous? <a
        href="https://docs.google.com/forms/d/1NI2nuAoYRZzYcawLrbWPKHsc43Edv"
        target="_blank">Let us know!</a>
</div>
<div ng-switch-when="disconnected">
    You have been disconnected.
</div>
<div ng-switch-when="error">

```

```

Sorry, we're unable to connect you. Please check the
following:
<ol>
  <li>
    You need to be using a computer connected to Princeton's
    network.<br>
    If you're off-campus, <a href="about#offcampus">follow
    these instructions.</a>
  </li>
  <li>You can't already be chatting with a user.</li>
  <li>You need to be using a modern web browser that
    supports WebSockets.</li>
</ol>
</div>
<div ng-switch-default>
  {{message.text}}
</div>
</div>
</div>
</li>
<li class="typing" ng-show="partnerTyping && state == 'chatting'">
  {{partnerName}} is typing...
</ul>
</div>
<div class="input-wrapper">
  <textarea
    tabindex="1"
    pom-focus-on-chat
    ng-disabled="state != 'chatting'"
    ng-model="message"
    ng-keydown="sendMessage($event)"
    ng-change="updateTyping()"></textarea>
</div>
</div>
<audio pom-play-on-message src="audio/notification.wav"></audio>
<script src="/socket.io.js"></script>
<script
  src="//ajax.googleapis.com/ajax/libs/angularjs/1.2.6/angular.min.js"></script>
<script
  src="//ajax.googleapis.com/ajax/libs/angularjs/1.2.6/angular-sanitize.js"></script>
<script
  src="//ajax.googleapis.com/ajax/libs/angularjs/1.2.6/angular-animate.js"></script>
<!-- build:js js/app.js -->
<script src="js/app.js"></script>
<script src="js/controllers.js"></script>
<script src="js/directives.js"></script>
<script src="js/services.js"></script>

```

```
<script src="js/filters.js"></script>
<!-- endbuild -->
</body>
</html>
```

Bibliography

Peter Auer, Nicoló Cesa-Bianchi, and Paul Fischer. Finite-time analysis of the multiarmed bandit problem. *Machine Learning*, 47:235–256, 2002.

Sébastien Bubeck and Nicoló Cesa-Bianchi. Regret analysis of stochastic and non-stochastic multi-armed bandit problems. *Foundations and Trends in Machine Learning*, 5(1):1–122, 2012.

Lihong Li, Wei Chu, John Langford, and Robert E. Schapire. A contextual-bandit approach to personalized news article recommendation. *Proceedings of the 19th International Conference on World Wide Web*, pages 661–670, 2010.